Dynamic Data >
Main Challenges

- Fast boot
- Garbage collection
- Persistence
- Replication performance
- Network efficiency

Fast Boot / catchup

For a new replica entering the network, replicating the entire history for CmRDTs is not scalable.
- Use consistent snapshotting? How?
- Use CvRDTs instead (and loose implicit op history and authorship)?